

In [4]:

```
1 import pandas as pd
2 data = pd.read_csv(r"C:\Users\manju\Desktop\HousingData.csv")
3 data
4 df=pd.DataFrame(data)
5 df
```

Out[4]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	L
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	
...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	21.0	396.90	

506 rows × 14 columns



## Boston house data set

### Data splitting

In [4]:

```
1 import pandas as pd
2 data = pd.read_csv(r"C:\Users\manju\Desktop\HousingData.csv")
3
```

```
In [6]: 1 X = data.drop('AGE', axis=1)
2 y = data['ZN']
3 from sklearn.model_selection import train_test_split
4 X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3,
5 X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_si
6 print(X)
7 print(y)
8
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	6.0622	3	222	18.7	
..	...	...	...	...	...	...	...	...	...	...	
501	0.06263	0.0	11.93	0.0	0.573	6.593	2.4786	1	273	21.0	
502	0.04527	0.0	11.93	0.0	0.573	6.120	2.2875	1	273	21.0	
503	0.06076	0.0	11.93	0.0	0.573	6.976	2.1675	1	273	21.0	
504	0.10959	0.0	11.93	0.0	0.573	6.794	2.3889	1	273	21.0	
505	0.04741	0.0	11.93	0.0	0.573	6.030	2.5050	1	273	21.0	

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2
..	...	...	...
501	391.99	NaN	22.4
502	396.90	9.08	20.6
503	396.90	5.64	23.9
504	393.45	6.48	22.0
505	396.90	7.88	11.9

[506 rows x 13 columns]

0	18.0
1	0.0
2	0.0
3	0.0
4	0.0

..	...
501	0.0
502	0.0
503	0.0
504	0.0
505	0.0

Name: ZN, Length: 506, dtype: float64

## train dataset

```
In [7]: 1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 data = pd.read_csv(r"C:\Users\manju\Desktop\HousingData.csv")
4 X = data.drop('AGE', axis=1)
5 y = data['TAX']
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
7 X
8
9
```

```
Out[7]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	4.0900	1	296	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	4.9671	2	242	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	4.9671	2	242	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	6.0622	3	222	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	6.0622	3	222	18.7	396.90	NaN
...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.0	0.573	6.593	2.4786	1	273	21.0	391.99	NaN
502	0.04527	0.0	11.93	0.0	0.573	6.120	2.2875	1	273	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	2.1675	1	273	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	2.3889	1	273	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	2.5050	1	273	21.0	396.90	7.88

506 rows × 13 columns



## Validation dataset


```
In [11]: 1 import csv
2 from sklearn.model_selection import train_test_split
3 data = []
4 with open(r"C:\Users\manju\Desktop\HousingData.csv") as file:
5     csv_reader = csv.reader(file)
6     for row in csv_reader:
7         data.append(row)
8 X = [row[:-1] for row in data]
9 y = [row[-1] for row in data]
10 X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
11 X
```

```
Out[11]: [['CRIM',
'ZN',
'INDUS',
'CHAS',
'NOX',
'RM',
'AGE',
'DIS',
'RAD',
'TAX',
'PTRATIO',
'B',
'LSTAT'],
['0.00632',
'18',
'2.31',
'0',
'0.538',
'6.575',
'0.0918']]
```

## Overfitting

```
In [14]: 1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.tree import DecisionTreeRegressor
4 from sklearn.metrics import mean_squared_error
5 import matplotlib.pyplot as plt
6 data = pd.read_csv(r"C:\Users\manju\Desktop\HousingData.csv")
7 X = data.drop('AGE', axis=1)
8 y = data['TAX']
9 X.fillna(X.mean(), inplace=True)
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
11 model = DecisionTreeRegressor(max_depth=None)
12 model.fit(X_train, y_train)
13 y_pred = model.predict(X_test)
14 mse = mean_squared_error(y_test, y_pred)
15 print(f"Mean Squared Error: {mse}")
16 plt.scatter(y_test, y_pred)
17 plt.xlabel("Actual Values")
18 plt.ylabel("Predicted Values")
19 plt.title("Actual vs. Predicted Values")
20 plt.show()
21
```

Mean Squared Error: 1.2843137254901962




## Underfitting

```

In [2]: 1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.linear_model import LinearRegression
4 from sklearn.metrics import mean_squared_error
5 import matplotlib.pyplot as plt
6
7 data = pd.read_csv(r"C:\Users\manju\Desktop\HousingData.csv")
8 X = data.drop('AGE', axis=1)
9 y = data['TAX']
10 X.fillna(X.mean(), inplace=True)
11 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
12 model = LinearRegression()
13 model.fit(X_train, y_train)
14 y_pred = model.predict(X_test)
15 mse = mean_squared_error(y_test, y_pred)
16 print(f"Mean Squared Error: {mse}")
17 plt.scatter(y_test, y_pred)
18 plt.xlabel("Actual Values")
19 plt.ylabel("Predicted Values")
20 plt.title("Actual vs. Predicted Values")
21 plt.show()
22

```

Mean Squared Error: 6.842486684721736e-27




**simple linear regression**

```

In [10]: 1 import pandas as pd
2 import numpy as np
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LinearRegression
5 import matplotlib.pyplot as plt
6 data = pd.read_csv(r"C:\Users\manju\Desktop\HousingData.csv")
7 X = data[['NOX']]
8 y = data['TAX']
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
10 model = LinearRegression()
11 model.fit(X_train, y_train)
12 y_pred = model.predict(X_test)
13 plt.scatter(X, y, color='blue', label='Data')
14 plt.plot(X_test, y_pred, color='red', linewidth=2, label='Regression Line')
15 plt.xlabel('Independent Variable')
16 plt.ylabel('Dependent Variable')
17 plt.legend()
18 plt.show()
19 slope = model.coef_[0]
20 intercept = model.intercept_
21 print(f"Slope (m): {slope}")
22 print(f"Intercept (b): {intercept}")
23

```



Slope (m): 1027.9020979428083  
Intercept (b): -168.55195983014949

## Multiple linear regression

```
In [13]: 1 import pandas as pd
2 import numpy as np
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LinearRegression
5 from sklearn.metrics import mean_squared_error, r2_score
6 data = pd.read_csv(r"C:\Users\manju\Desktop\HousingData.csv")
7 print(data.head())
8 print(data.info())
9 X = data[['NOX']]
10 y = data['TAX']
11 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
12 model = LinearRegression()
13 model.fit(X_train, y_train)
14 y_pred = model.predict(X_test)
15 mse = mean_squared_error(y_test, y_pred)
16 r2 = r2_score(y_test, y_pred)
17 print("Mean Squared Error:", mse)
18 print("R-squared:", r2)
19 print("Coefficients:", model.coef_)
20 print("Intercept:", model.intercept_)
```



	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATI
0 \	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.
3											
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.
8											
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.
8											
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.
7											
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.
7											

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    CRIM        486 non-null    float64
1    ZN          486 non-null    float64
2    INDUS       486 non-null    float64
3    CHAS        486 non-null    float64
4    NOX         506 non-null    float64
5    RM          506 non-null    float64
6    AGE         486 non-null    float64
7    DIS         506 non-null    float64
8    RAD         506 non-null    int64
9    TAX         506 non-null    int64
10   PTRATIO     506 non-null    float64
11   B           506 non-null    float64
12   LSTAT       486 non-null    float64
13   MEDV       506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
None
Mean Squared Error: 15264.37775481554
R-squared: 0.510508842100377
Coefficients: [918.00440183]
Intercept: -106.82272875599114

```

## polynomial linear regression

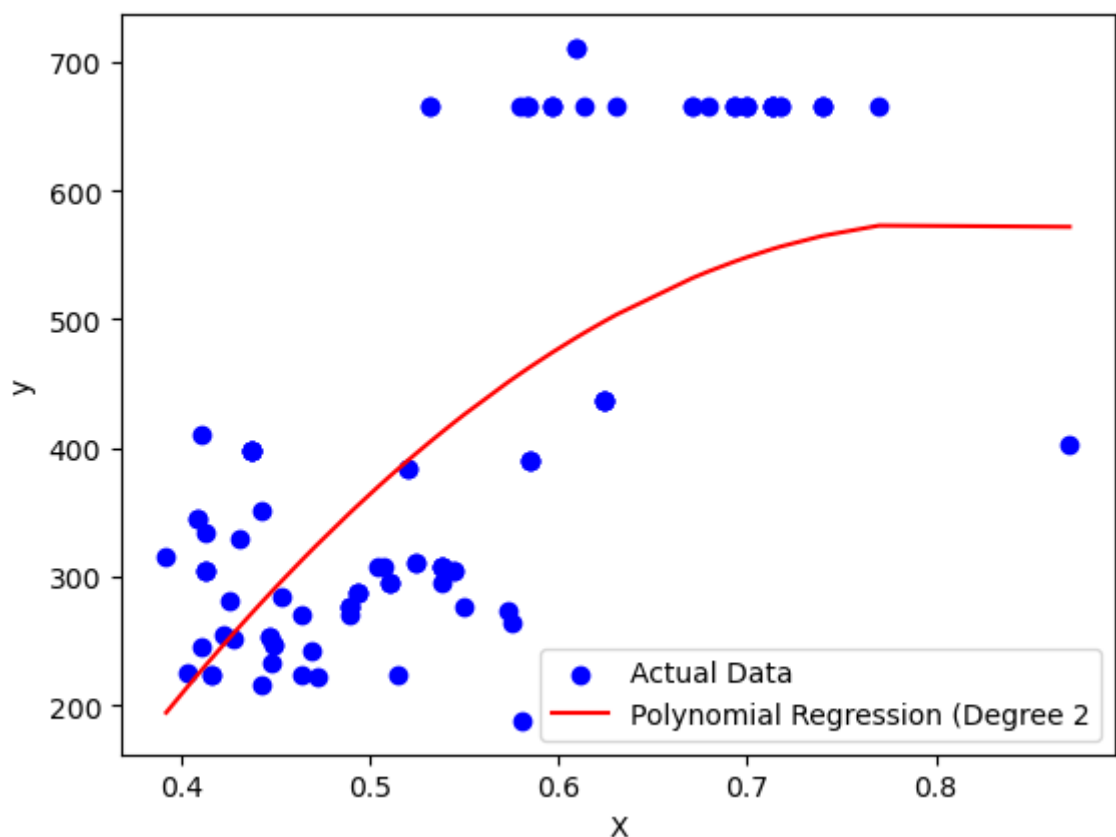
```

In [39]: 1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LinearRegression
5 from sklearn.preprocessing import PolynomialFeatures
6 from sklearn.metrics import mean_squared_error, r2_score
7 import matplotlib.pyplot as plt
8 X = data[['NOX']]
9 y = data['TAX']
10 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
11 degree = 2
12 poly_features = PolynomialFeatures(degree=degree)
13 X_train_poly = poly_features.fit_transform(X_train)
14 X_test_poly = poly_features.transform(X_test)
15 poly_reg = LinearRegression()
16 poly_reg.fit(X_train_poly, y_train)
17 y_pred = poly_reg.predict(X_test_poly)
18 mse = mean_squared_error(y_test, y_pred)
19 r2 = r2_score(y_test, y_pred)
20 print("Mean Squared Error:", mse)
21 print("R-squared (R2) Score:", r2)
22 X_test_sorted, y_pred_sorted = zip(*sorted(zip(X_test.values, y_pred)))
23 plt.scatter(X_test, y_test, color='blue', label='Actual Data')
24 plt.plot(X_test_sorted, y_pred_sorted, color='red', label=f'Polynomial Regression (Degree 2)')
25 plt.xlabel('X')
26 plt.ylabel('y')
27 plt.legend()
28 plt.show()

```


Mean Squared Error: 14546.715241982676

R-squared (R2) Score: 0.5335225187814907



# MAE- mean absolute Error

```
In [7]: 1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.linear_model import LinearRegression
4 from sklearn.metrics import mean_absolute_error
5 np.random.seed(0)
6 X = np.random.rand(100, 1) * 10
7 y = 2 * X + 1 + np.random.randn(100, 1)
8 X_train, X_test, y_train, y_test = X[:80], X[80:], y[:80], y[80:]
9 model = LinearRegression()
10 model.fit(X_train, y_train)
11 predictions = model.predict(X_test)
12 mae = mean_absolute_error(y_test, predictions)
13 plt.figure(figsize=(8, 6))
14 plt.scatter(X_test, y_test, color='blue', label='Actual')
15 plt.scatter(X_test, predictions, color='red', marker='x', label='Predicted')
16 plt.plot(X_test, predictions, color='green', linewidth=2, label='Regression')
17 plt.xlabel('X')
18 plt.ylabel('y')
19 plt.title(f'Linear Regression (MAE: {mae:.2f})')
20 plt.legend()
21 plt.grid()
22 plt.show()
23
```




# MSE- mean squared error

```

In [18]: 1
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.linear_model import LinearRegression
5 from sklearn.metrics import mean_squared_error
6 np.random.seed(0)
7 X = np.random.rand(100, 1) * 10
8 y = 2 * X + 1 + np.random.randn(100, 1)
9 X_train, X_test, y_train, y_test = X[:80], X[80:], y[:80], y[80:]
10 model = LinearRegression()
11 model.fit(X_train, y_train)
12 predictions = model.predict(X_test)
13 squared_errors = (predictions - y_test) ** 2
14 rmse = np.sqrt(mean_squared_error(y_test, predictions))
15 plt.figure(figsize=(8, 6))
16 plt.scatter(X_test, y_test, color='blue', label='Actual')
17 plt.scatter(X_test, predictions, color='red', marker='x', label='Predictions')
18 plt.plot(X_test, predictions, color='green', linewidth=2, label='Regression Line')
19 plt.xlabel('X')
20 plt.ylabel('Y')
21 plt.title(f'Linear Regression (RMSE: {rmse:.2f})')

```

Out[18]: Text(0.5, 1.0, 'Linear Regression (RMSE: 0.95)')



```
In [22]: 1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.linear_model import LinearRegression
4 from sklearn.metrics import mean_squared_error
5 np.random.seed(0)
6 X = np.random.rand(100, 1) * 10
7 y = 2 * X + 1 + np.random.randn(100, 1)
8 X_train, X_test, y_train, y_test = X[:80], X[80:], y[:80], y[80:]
9 model = LinearRegression()
10 model.fit(X_train, y_train)
11 predictions = model.predict(X_test)
12 mae = mean_absolute_error(y_test, predictions)
13 plt.figure(figsize=(8, 6))
14 plt.scatter(X_test, y_test, color='blue', label='Actual')
15 plt.scatter(X_test, predictions, color='red', marker='x', label='Predictions')
16 plt.plot(X_test, predictions, color='green', linewidth=2, label='Regression Line')
17 plt.xlabel('X')
18 plt.ylabel('y')
19 plt.title(f'Linear Regression (MSE: {mse:.2f})')
20 plt.legend()
21 plt.grid()
22 plt.show()
23
```

-----  
**ImportError** Traceback (most recent call last)

Cell In[22], line 4

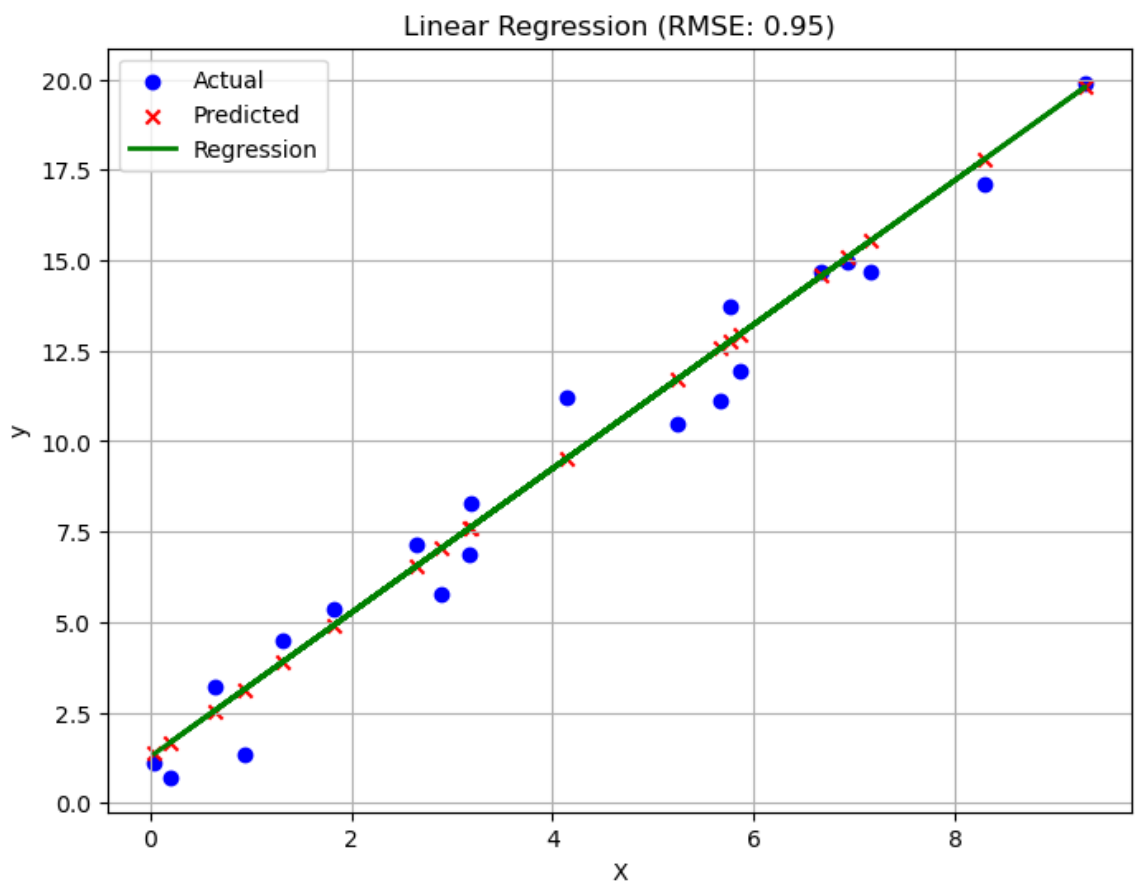
```
2 import matplotlib.pyplot as plt
3 from sklearn.linear_model import LinearRegression
----> 4 from sklearn.metrics import mean_squared_error
5 np.random.seed(0)
6 X = np.random.rand(100, 1) * 10
```

**ImportError:** cannot import name 'mean\_squared\_error' from 'sklearn.metrics' (C:\Users\manju\anaconda3\Lib\site-packages\sklearn\metrics\\_\_init\_\_.py)

```

In [20]: 1 #root mean squared error
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.linear_model import LinearRegression
5 from sklearn.metrics import mean_squared_error
6 np.random.seed(0)
7 X = np.random.rand(100, 1) * 10
8 y = 2 * X + 1 + np.random.randn(100, 1)
9 X_train, X_test, y_train, y_test = X[:80], X[80:], y[:80], y[80:]
10 model = LinearRegression()
11 model.fit(X_train, y_train)
12 predictions = model.predict(X_test)
13 squared_errors = (predictions - y_test) ** 2
14 rmse = np.sqrt(mean_squared_error(y_test, predictions))
15 plt.figure(figsize=(8, 6))
16 plt.scatter(X_test, y_test, color='blue', label='Actual')
17 plt.scatter(X_test, predictions, color='red', marker='x', label='Predicted')
18 plt.plot(X_test, predictions, color='green', linewidth=2, label='Regression')
19 plt.xlabel('X')
20 plt.ylabel('y')
21 plt.title(f'Linear Regression (RMSE: {rmse:.2f})')
22 plt.legend()
23 plt.grid()
24 plt.show()
25
26

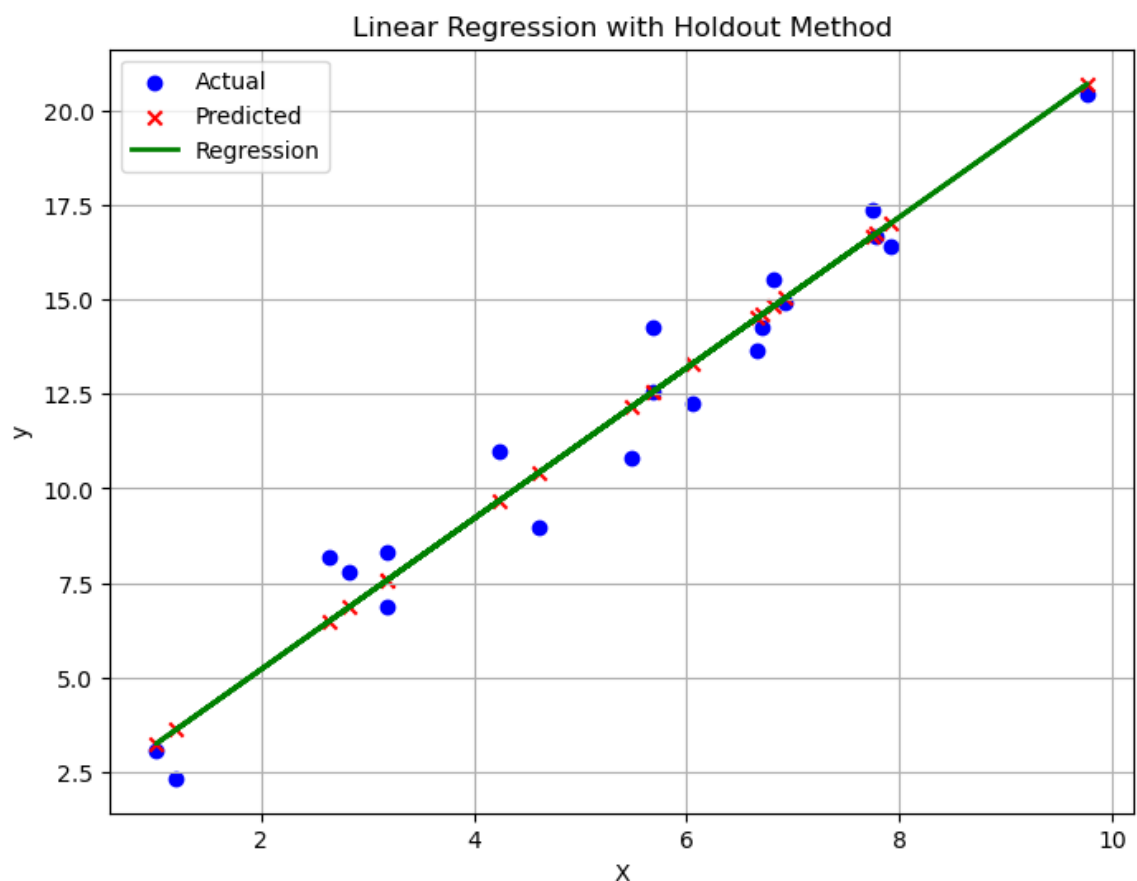
```



**Hold out method**

In [25]:

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LinearRegression
5 np.random.seed(0)
6 X = np.random.rand(100, 1) * 10
7 y = 2 * X + 1 + np.random.randn(100, 1)
8 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
9 model = LinearRegression()
10 model.fit(X_train, y_train)
11 predictions = model.predict(X_test)
12 plt.figure(figsize=(8, 6))
13 plt.scatter(X_test, y_test, color='blue', label='Actual')
14 plt.scatter(X_test, predictions, color='red', marker='x', label='Predicted')
15 plt.plot(X_test, predictions, color='green', linewidth=2, label='Regression')
16 plt.xlabel('X')
17 plt.ylabel('y')
18 plt.title('Linear Regression with Holdout Method')
19 plt.legend()
20 plt.grid()
21 plt.show()
22
23
```



In [29]:

```
1 # Loading the breast_cancer dataset and finding the accuracy and classif
```

In [28]:

```
1
2 from sklearn.datasets import load_breast_cancer
3 from sklearn.model_selection import train_test_split
4 from sklearn.tree import DecisionTreeClassifier
5 from sklearn.metrics import accuracy_score, classification_report
6 breast_cancer = load_breast_cancer()
7 X = breast_cancer.data # Features
8 y = breast_cancer.target # Target variable
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
10 clf = DecisionTreeClassifier(random_state=42)
11 clf.fit(X_train, y_train)
12 y_pred = clf.predict(X_test)
13 accuracy = accuracy_score(y_test, y_pred)
14 print(f"Accuracy: {accuracy:.2f}")
15 print("Classification Report:")
16 print(classification_report(y_test, y_pred))
17
```

Accuracy: 0.95

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.93	0.93	43
1	0.96	0.96	0.96	71
accuracy			0.95	114
macro avg	0.94	0.94	0.94	114
weighted avg	0.95	0.95	0.95	114

## r-square



In [1]:

```
1 #R-SQUARE
2 import numpy as np
3 from sklearn.linear_model import LinearRegression
4 from sklearn.metrics import r2_score
5 np.random.seed(0)
6 X = np.random.rand(100, 1) * 10
7 y = 2 * X + 1 + np.random.randn(100, 1)
8 X_train, X_test, y_train, y_test = X[:80], X[80:], y[:80], y[80:]
9 model = LinearRegression()
10 model.fit(X_train, y_train)
11 predictions = model.predict(X_test)
12 r_squared = r2_score(y_test, predictions)
13 print("X_train:", X_train)
14 print("y_train:", y_train)
15 print("X_test:", X_test)
16 print("y_test:", y_test)
17 print(f'R-squared: {r_squared:.2f}')
18
```

X\_train: [[5.48813504]

[7.15189366]

[6.02763376]

[5.44883183]

[4.23654799]

[6.45894113]

[4.37587211]

[8.91773001]

[9.63662761]

[3.83441519]

[7.91725038]

[5.2889492 ]

[5.68044561]

[9.25596638]

[0.71036058]

[0.871293 ]

[0.20218397]

[8.32619846]

[7.78156751]

5.07001001403

```

In [8]: 1 import numpy as np
        2 import matplotlib.pyplot as plt
        3 n_train = 150
        4 n_test = 1000
        5 noise = 0.1
        6 def f(x):
        7     x = x.ravel()
        8     return np.exp(-x ** 2) + 1.5 * np.exp(-(x - 2) ** 2)
        9 def generate(n_samples, noise):
       10     X = np.random.rand(n_samples) * 10 - 5
       11     X = np.sort(X).ravel()
       12     y = np.exp(-X ** 2) + 1.5 * np.exp(-(X - 2) ** 2) + np.random.normal(0
       13     X = X.reshape((n_samples, 1))
       14     return X, y
       15 X_train, y_train = generate(n_samples=n_train, noise=noise)
       16 X_test, y_test = generate(n_samples=n_test, noise=noise)

```

```


In [ ]: 1

```

```

In [7]: 1 plt.figure(figsize=(18, 8))
        2 plt.plot(X_test, f(X_test), "r")
        3 plt.scatter(X_train, y_train, c="b", s=20)
        4 plt.xlim([-5, 5])
        5 plt.show()
        6

```



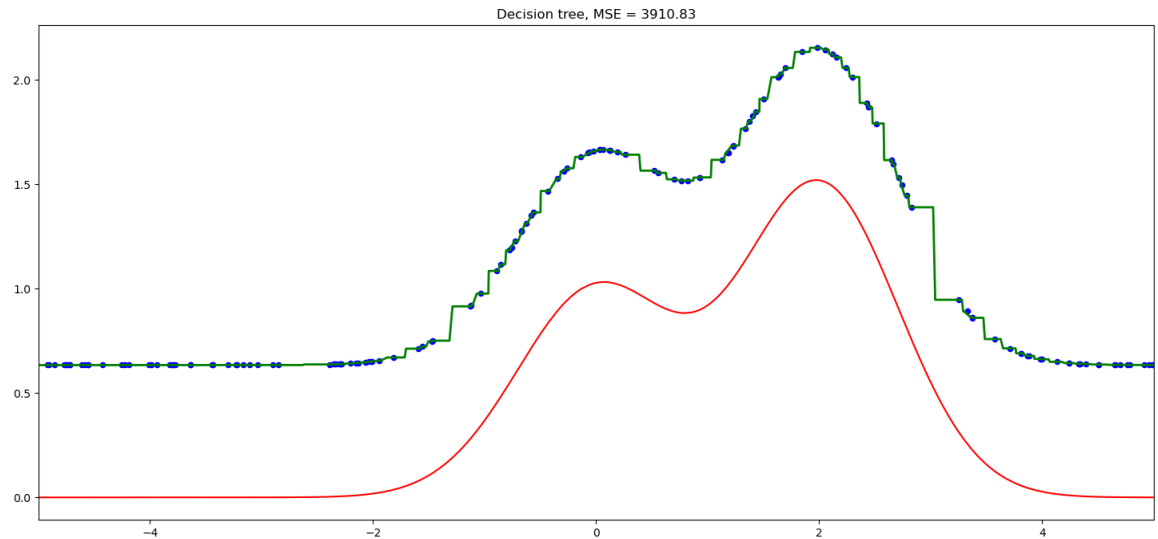
## decision tree Regresson

```

In [9]: 1 #DecisionTreeRegressor
2 from sklearn.tree import DecisionTreeRegressor
3 dtree = DecisionTreeRegressor().fit(X_train, y_train)
4 d_predict = dtree.predict(X_test)
5 plt.figure(figsize=(18, 8))
6 plt.plot(X_test, f(X_test), "r")
7 plt.scatter(X_train, y_train, c="b", s=20)
8 plt.plot(X_test, d_predict, "g", lw=2)
9 plt.xlim([-5, 5])
10 plt.title("Decision tree, MSE = %.2f"
11 % np.sum((y_test - d_predict) ** 2))

```

Out[9]: Text(0.5, 1.0, 'Decision tree, MSE = 3910.83')



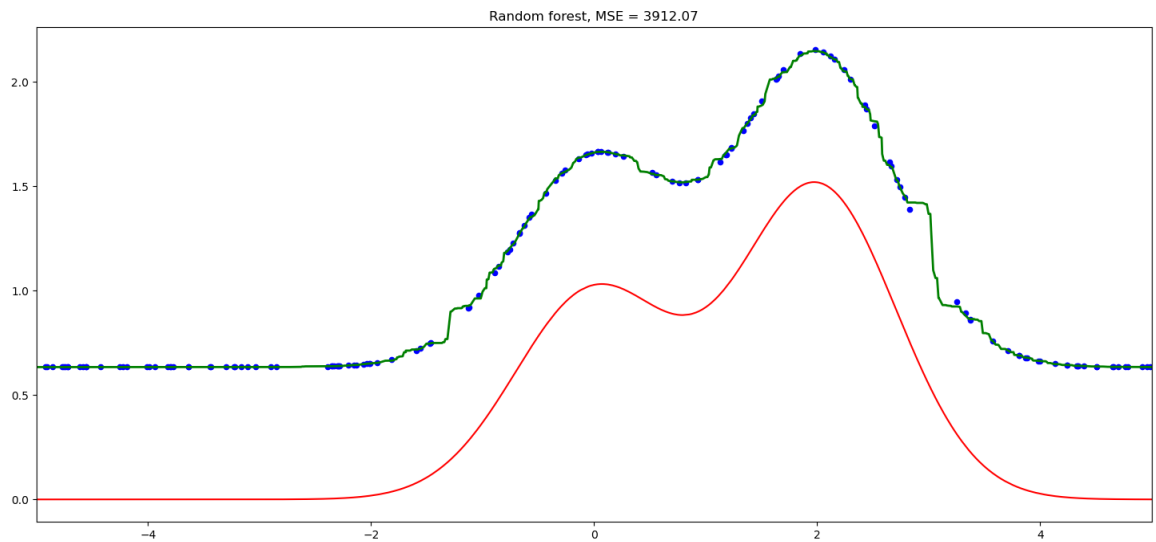
## random forest algorithm

```

In [11]: 1 from sklearn.ensemble import RandomForestRegressor
2 rfr = RandomForestRegressor(n_estimators=1000).fit(X_train, y_train)
3 rf_predict = rfr.predict(X_test)
4 plt.figure(figsize=(18, 8))
5 plt.plot(X_test, f(X_test), "r")
6 plt.scatter(X_train, y_train, c="b", s=20)
7 plt.plot(X_test, rf_predict, "g", lw=2)
8 plt.xlim([-5, 5])
9 plt.title("Random forest, MSE = %.2f" % np.sum((y_test - rf_predict) **
10

```

Out[11]: Text(0.5, 1.0, 'Random forest, MSE = 3912.07')



In [ ]: 1